Synopsis

An example of a web server written with Node which responds with 'Hello World':

```
var http = require('http');

http.createServer(function (request, response) {
  response.writeHead(200, {'Content-Type': 'text/plain'});
  response.end('Hello World\n');
}).listen(8124);

console.log('Server running at http://127.0.0.1:8124/');
```

To run the server, put the code into a file called `example.js` and execute it with the node program

```
> node example.js
Server running at http://127.0.0.1:8124/
```

All of the examples in the documentation can be run similarly.

Standard Modules

Node comes with a number of modules that are compiled in to the process, most of which are documented below. The most common way to use these modules is with `require('name')` and then assigning the return value to a local variable with the same name as the module.

Example:

```
var sys = require('sys');
```

It is possible to extend node with other modules. See `'Modules'`

Buffers

Pure Javascript is Unicode friendly but not nice to binary data. When dealing with TCP streams or the file system, it's necessary to handle octet streams. Node has several strategies for manipulating, creating, and consuming octet streams.

Raw data is stored in instances of the `Buffer` class. A `Buffer` is similar to an array of integers but corresponds to a raw memory allocation outside the V8 heap. A `Buffer` cannot be resized.

The `Buffer` object is global.

Converting between Buffers and JavaScript string objects requires an explicit encoding method. Here are the different string encodings;

- `'ascii'` - for 7 bit ASCII data only. This encoding method is very fast, and will strip the high bit if set.

- `'utf8'` - Unicode characters. Many web pages and other document formats use UTF-8.

- `'base64'` - Base64 string encoding.

- `'binary'` - A way of encoding raw binary data into strings by using only the first 8 bits of each character. This encoding method is depreciated and should be avoided in favor of `Buffer` objects where possible. This encoding will be removed in future versions of Node.

new Buffer(size)
Allocates a new buffer of `size` octets.

new Buffer(array)
Allocates a new buffer using an `array` of octets.

new Buffer(str, encoding='utf8')
Allocates a new buffer containing the given `str`.

buffer.write(string, offset=0, encoding='utf8')
Writes `string` to the buffer at `offset` using the given encoding. Returns number of octets written. If `buffer` did not contain enough space to fit the entire string it will write a partial amount of the string. In the case of `'utf8'` encoding, the method will not write partial characters.

Example: write a utf8 string into a buffer, then print it

```
buf = new Buffer(256);
len = buf.write('\u00bd + \u00bc = \u00be', 0);
console.log(len + " bytes: " + buf.toString('utf8', 0, len));

// 12 bytes: ½ + ¼ = ¾
```

buffer.toString(encoding, start=0, end=buffer.length)
Decodes and returns a string from buffer data encoded with `encoding` beginning at `start` and ending at `end`.

See `buffer.write()` example, above.

buffer[index]
Get and set the octet at `index`. The values refer to individual bytes, so the legal range is between `0x00` and `0xFF` hex or `0` and `255`.

Example: copy an ASCII string into a buffer, one byte at a time:

```
str = "node.js";
buf = new Buffer(str.length);

for (var i = 0; i < str.length ; i++) {
  buf[i] = str.charCodeAt(i);
}

console.log(buf);

// node.js
```

Buffer.byteLength(string, encoding='utf8')

Gives the actual byte length of a string. This is not the same as `string.prototype.length` since that returns the number of <u>characters</u> in a string.

Example:

```
str = '\u00bd + \u00bc = \u00be';

console.log(str + ": " + str.length + " characters, " +
  Buffer.byteLength(str, 'utf8') + " bytes");

// ½ + ¼ = ¾: 9 characters, 12 bytes
```

buffer.length

The size of the buffer in bytes. Note that this is not necessarily the size of the contents. `length` refers to the amount of memory allocated for the buffer object. It does not change when the contents of the buffer are changed.

```
buf = new Buffer(1234);

console.log(buf.length);
buf.write("some string", "ascii", 0);
console.log(buf.length);

// 1234
// 1234
```

buffer.copy(targetBuffer, targetStart, sourceStart, sourceEnd=buffer.length)

Does a memcpy() between buffers.

Example: build two Buffers, then copy `buf1` from byte 16 through byte 19 into `buf2`, starting at the 8th byte in `buf2`.

```
buf1 = new Buffer(26);
```

```
  buf2 = new Buffer(26);

  for (var i = 0 ; i < 26 ; i++) {
    buf1[i] = i + 97; // 97 is ASCII a
    buf2[i] = 33; // ASCII !
  }

  buf1.copy(buf2, 8, 16, 20);
  console.log(buf2.toString('ascii', 0, 25));


  // !!!!!!!!qrst!!!!!!!!!!!!!!
```

buffer.slice(start, end=buffer.length)

Returns a new buffer which references the same memory as the old, but offset and cropped by the `start` and `end` indexes.

Modifying the new buffer slice will modify memory in the original buffer!

Example: build a Buffer with the ASCII alphabet, take a slice, then modify one byte from the original Buffer.

```
var buf1 = new Buffer(26);

for (var i = 0 ; i < 26 ; i++) {
  buf1[i] = i + 97; // 97 is ASCII a
}

var buf2 = buf1.slice(0, 3);
console.log(buf2.toString('ascii', 0, buf2.length));
buf1[0] = 33;
console.log(buf2.toString('ascii', 0, buf2.length));

// abc
// !bc
```

EventEmitter

Many objects in Node emit events: a TCP server emits an event each time there is a stream, a child process emits an event when it exits. All objects which emit events are instances of `events.EventEmitter`.

Events are represented by a camel-cased string. Here are some examples: `'stream'`, `'data'`, `'messageBegin'`.

Functions can be then be attached to objects, to be executed when an event is emitted. These functions are called listeners.

`require('events').EventEmitter` to access the `EventEmitter` class.

All EventEmitters emit the event `'newListener'` when new listeners are added.

When an `EventEmitter` experiences an error, the typical action is to emit an `'error'` event. Error events are special--if there is no handler for them they will print a stack trace and exit the program.

Event: 'newListener'

```
function (event, listener) { }
```

This event is emitted any time someone adds a new listener.

Event: 'error'

```
function (exception) { }
```

If an error was encountered, then this event is emitted. This event is special - when there are no listeners to receive the error Node will terminate execution and display the exception's stack trace.

emitter.on(event, listener)

Adds a listener to the end of the listeners array for the specified event.

```
server.on('stream', function (stream) {
  console.log('someone connected!');
});
```

emitter.removeListener(event, listener)

Remove a listener from the listener array for the specified event. **Caution**: changes array indices in the listener array behind the listener.

```
var callback = function(stream) {
  console.log('someone connected!');
};
server.on('stream', callback);
// ...
server.removeListener('stream', callback);
```

emitter.removeAllListeners(event)

Removes all listeners from the listener array for the specified event.

emitter.listeners(event)

Returns an array of listeners for the specified event. This array can be manipulated, e.g. to remove listeners.

```
server.on('stream', function (stream) {
  console.log('someone connected!');
});
console.log(sys.inspect(server.listeners('stream'));
```

```
// [ [Function] ]
```

emitter.emit(event, [arg1], [arg2], [...])
Execute each of the listeners in order with the supplied arguments.

A stream is an abstract interface implemented by various objects in Node. For example a request to an HTTP server is a stream, as is stdout. Streams are readable, writable, or both. All streams are instances of `EventEmitter`.

Readable Stream

A `Readable Stream` has the following methods, members, and events.

Event: 'data'
```
function (data) { }
```

The `'data'` event emits either a `Buffer` (by default) or a string if `setEncoding()` was used.

Event: 'end'
```
function () { }
```

Emitted when the stream has received an EOF (FIN in TCP terminology). Indicates that no more `'data'` events will happen. If the stream is also writable, it may be possible to continue writing.

Event: 'error'
```
function (exception) { }
```

Emitted if there was an error receiving data.

Event: 'close'
```
function () { }
```

Emitted when the underlying file descriptor has be closed. Not all streams will emit this. (For example, an incoming HTTP request will not emit `'close'`.)

Event: 'fd'
```
function (fd) { }
```

Emitted when a file descriptor is received on the stream. Only UNIX streams support this functionality; all others will simply never emit this event.

stream.readable
A boolean that is `true` by default, but turns `false` after an `'error'` occured, the stream came to an `'end'`, or `destroy()` was called.

stream.setEncoding(encoding)

Makes the data event emit a string instead of a `Buffer`. `encoding` can be `'utf8'`, `'ascii'`, or `'base64'`.

stream.pause()

Pauses the incoming `'data'` events.

stream.resume()

Resumes the incoming `'data'` events after a `pause()`.

stream.destroy()

Closes the underlying file descriptor. Stream will not emit any more events.

Writable Stream

A `Writable Stream` has the following methods, members, and events.

Event: 'drain'

```
function () { }
```

Emitted after a `write()` method was called that returned `false` to indicate that it is safe to write again.

Event: 'error'

```
function (exception) { }
```

Emitted on error with the exception `exception`.

Event: 'close'

```
function () { }
```

Emitted when the underlying file descriptor has been closed.

stream.writeable

A boolean that is `true` by default, but turns `false` after an `'error'` occurred or `end()` / `destroy()` was called.

stream.write(string, encoding='utf8', [fd])

Writes `string` with the given `encoding` to the stream. Returns `true` if the string has been flushed to the kernel buffer. Returns `false` to indicate that the kernel buffer is full, and the data will be sent out in the future. The `'drain'` event will indicate when the kernel buffer is empty again. The `encoding` defaults to `'utf8'`.

If the optional `fd` parameter is specified, it is interpreted as an integral file descriptor to be sent over the stream. This is only supported for UNIX streams, and is silently ignored otherwise. When writing a file descriptor in this manner, closing the descriptor before the stream drains risks sending an invalid (closed) FD.

stream.write(buffer)

Same as the above except with a raw buffer.

stream.end()

Terminates the stream with EOF or FIN.

stream.end(string, encoding)

Sends `string` with the given `encoding` and terminates the stream with EOF or FIN. This is useful to reduce the number of packets sent.

stream.end(buffer)

Same as above but with a `buffer`.

stream.destroy()

Closes the underlying file descriptor. Stream will not emit any more events.

Global Objects

These object are available in the global scope and can be accessed from anywhere.

global

The global namespace object.

process

The process object. See the `'process object'` section.

require()

To require modules. See the `'Modules'` section.

require.paths

An array of search paths for `require()`. This array can be modified to add custom paths.

Example: add a new path to the beginning of the search list

```
require.paths.unshift('/usr/local/node');
console.log(require.paths);
// /usr/local/node,/Users/mjr/.node_libraries
```

__filename

The filename of the script being executed. This is the absolute path, and not necessarily the same filename passed in as a command line argument.

Example: running `node example.js` from `/Users/mjr`

```
console.log(__filename);
// /Users/mjr/example.js
```

__dirname

The dirname of the script being executed.

Example: running `node example.js` from `/Users/mjr`

```
console.log(__dirname);
// /Users/mjr
```

module

A reference to the current module (of type `process.Module`). In particular `module.exports` is the same as the `exports` object. See `src/process.js` for more information.

process

The `process` object is a global object and can be accessed from anywhere. It is an instance of `EventEmitter`.

Event: 'exit'
`function () {}`

Emitted when the process is about to exit. This is a good hook to perform constant time checks of the module's state (like for unit tests). The main event loop will no longer be run after the 'exit' callback finishes, so timers may not be scheduled.

Example of listening for `exit`:

```
process.on('exit', function () {
  process.nextTick(function () {
   console.log('This will not run');
  });
  console.log('About to exit.');
});
```

Event: 'uncaughtException'
`function (err) { }`

Emitted when an exception bubbles all the way back to the event loop. If a listener is added for this exception, the default action (which is to print a stack trace and exit) will not occur.

Example of listening for `uncaughtException`:

```
process.on('uncaughtException', function (err) {
  console.log('Caught exception: ' + err);
});
```

```
  setTimeout(function () {
    console.log('This will still run.');
  }, 500);

  // Intentionally cause an exception, but don't catch it.
  nonexistentFunc();
  console.log('This will not run.');
```

Note that `uncaughtException` is a very crude mechanism for exception handling. Using try / catch in your program will give you more control over your program's flow. Especially for server programs that are designed to stay running forever, `uncaughtException` can be a useful safety mechanism.

Signal Events
`function () {}`

Emitted when the processes receives a signal. See sigaction(2) for a list of standard POSIX signal names such as SIGINT, SIGUSR1, etc.

Example of listening for `SIGINT`:

```
var stdin = process.openStdin();

process.on('SIGINT', function () {
  console.log('Got SIGINT.  Press Control-D to exit.');
});
```

An easy way to send the `SIGINT` signal is with `Control-C` in most terminal programs.

process.stdout
A `Writable Stream` to `stdout`.

Example: the definition of `console.log`

```
console.log = function (d) {
  process.stdout.write(d + '\n');
};
```

process.openStdin()
Opens the standard input stream, returns a `Readable Stream`.

Example of opening standard input and listening for both events:

```
var stdin = process.openStdin();

stdin.setEncoding('utf8');

stdin.on('data', function (chunk) {
```

```
    process.stdout.write('data: ' + chunk);
  });

  stdin.on('end', function () {
    process.stdout.write('end');
  });
```

process.argv

An array containing the command line arguments. The first element will be 'node', the second element will be the name of the JavaScript file. The next elements will be any additional command line arguments.

```
// print process.argv
process.argv.forEach(function (val, index, array) {
  console.log(index + ': ' + val);
});
```

This will generate:

```
$ node process-2.js one two=three four
0: node
1: /Users/mjr/work/node/process-2.js
2: one
3: two=three
4: four
```

process.execPath

This is the absolute pathname of the executable that started the process.

Example:

```
/usr/local/bin/node
```

process.chdir(directory)

Changes the current working directory of the process or throws an exception if that fails.

```
console.log('Starting directory: ' + process.cwd());
try {
  process.chdir('/tmp');
  console.log('New directory: ' + process.cwd());
}
catch (err) {
  console.log('chdir: ' + err);
}
```

process.compile(code, filename)

Similar to `eval` except that you can specify a `filename` for better error reporting and the `code` cannot see the local scope. The value of `filename` will be used as a filename if a

stack trace is generated by the compiled code.

Example of using `process.compile` and `eval` to run the same code:

```
var localVar = 123,
    compiled, evaled;

compiled = process.compile('localVar = 1;', 'myfile.js');
console.log('localVar: ' + localVar + ', compiled: ' + compiled);
evaled = eval('localVar = 1;');
console.log('localVar: ' + localVar + ', evaled: ' + evaled);

// localVar: 123, compiled: 1
// localVar: 1, evaled: 1
```

`process.compile` does not have access to the local scope, so `localVar` is unchanged.
`eval` does have access to the local scope, so `localVar` is changed.

In case of syntax error in `code`, `process.compile` exits node.

See also: `Script`

process.cwd()

Returns the current working directory of the process.

```
console.log('Current directory: ' + process.cwd());
```

process.env

An object containing the user environment. See environ(7).

process.exit(code=0)

Ends the process with the specified `code`. If omitted, exit uses the 'success' code 0.

To exit with a 'failure' code:

```
process.exit(1);
```

The shell that executed node should see the exit code as 1.

process.getgid()

Gets the group identity of the process. (See getgid(2).) This is the numerical group id, not the group name.

```
console.log('Current gid: ' + process.getgid());
```

process.setgid(id)

Sets the group identity of the process. (See setgid(2).) This accepts either a numerical ID

or a groupname string. If a groupname is specified, this method blocks while resolving it to a numerical ID.

```
console.log('Current gid: ' + process.getgid());
try {
  process.setgid(501);
  console.log('New gid: ' + process.getgid());
}
catch (err) {
  console.log('Failed to set gid: ' + err);
}
```

process.getuid()

Gets the user identity of the process. (See getuid(2).) This is the numerical userid, not the username.

```
console.log('Current uid: ' + process.getuid());
```

process.setuid(id)

Sets the user identity of the process. (See setuid(2).) This accepts either a numerical ID or a username string. If a username is specified, this method blocks while resolving it to a numerical ID.

```
console.log('Current uid: ' + process.getuid());
try {
  process.setuid(501);
  console.log('New uid: ' + process.getuid());
}
catch (err) {
  console.log('Failed to set uid: ' + err);
}
```

process.version

A compiled-in property that exposes NODE_VERSION.

```
console.log('Version: ' + process.version);
```

process.installPrefix

A compiled-in property that exposes NODE_PREFIX.

```
console.log('Prefix: ' + process.installPrefix);
```

process.kill(pid, signal='SIGINT')

Send a signal to a process. pid is the process id and signal is the string describing the signal to send. Signal names are strings like 'SIGINT' or 'SIGUSR1'. If omitted, the signal will be 'SIGINT'. See kill(2) for more information.

Note that just because the name of this function is `process.kill`, it is really just a signal sender, like the `kill` system call. The signal sent may do something other than kill the target process.

Example of sending a signal to yourself:

```
process.on('SIGHUP', function () {
  console.log('Got SIGHUP signal.');
});

setTimeout(function () {
  console.log('Exiting.');
  process.exit(0);
}, 100);

process.kill(process.pid, 'SIGHUP');
```

process.pid
The PID of the process.

```
console.log('This process is pid ' + process.pid);
```

process.title
Getter/setter to set what is displayed in 'ps'.

process.platform
What platform you're running on. `'linux2'`, `'darwin'`, etc.

```
console.log('This platform is ' + process.platform);
```

process.memoryUsage()
Returns an object describing the memory usage of the Node process.

```
var sys = require('sys');

console.log(sys.inspect(process.memoryUsage()));
```

This will generate:

```
{ rss: 4935680
, vsize: 41893888
, heapTotal: 1826816
, heapUsed: 650472
}
```

`heapTotal` and `heapUsed` refer to V8's memory usage.

process.nextTick(callback)

On the next loop around the event loop call this callback. This is <u>not</u> a simple alias to `setTimeout(fn, 0)`, it's much more efficient.

```
process.nextTick(function () {
  console.log('nextTick callback');
});
```

process.umask([mask])

Sets or reads the process's file mode creation mask. Child processes inherit the mask from the parent process. Returns the old mask if `mask` argument is given, otherwise returns the current mask.

```
var oldmask, newmask = 0644;

oldmask = process.umask(newmask);
console.log('Changed umask from: ' + oldmask.toString(8) +
            ' to ' + newmask.toString(8));
```

sys

These functions are in the module `'sys'`. Use `require('sys')` to access them.

sys.print(string)

Like `console.log()` but without the trailing newline.

```
require('sys').print('String with no newline');
```

sys.debug(string)

A synchronous output function. Will block the process and output `string` immediately to `stderr`.

```
require('sys').debug('message on stderr');
```

sys.log(string)

Output with timestamp on `stdout`.

```
require('sys').log('Timestmaped message.');
```

sys.inspect(object, showHidden=false, depth=2)

Return a string representation of `object`, which is useful for debugging.

If `showHidden` is `true`, then the object's non-enumerable properties will be shown too.

If `depth` is provided, it tells `inspect` how many times to recurse while formatting the object. This is useful for inspecting large complicated objects.

The default is to only recurse twice. To make it recurse indefinitely, pass in `null` for `depth`.

Example of inspecting all properties of the `sys` object:

```
var sys = require('sys');

console.log(sys.inspect(sys, true, null));
```

sys.pump(readableStream, writeableStream, [callback])
Experimental

Read the data from `readableStream` and send it to the `writableStream`. When `writeableStream.write(data)` returns `false` `readableStream` will be paused until the `drain` event occurs on the `writableStream`. `callback` gets an error as its only argument and is called when `writableStream` is closed or when an error occurs.

Timers

setTimeout(callback, delay, [arg], [...])
To schedule execution of `callback` after `delay` milliseconds. Returns a `timeoutId` for possible use with `clearTimeout()`. Optionally, you can also pass arguments to the callback.

clearTimeout(timeoutId)
Prevents a timeout from triggering.

setInterval(callback, delay, [arg], [...])
To schedule the repeated execution of `callback` every `delay` milliseconds. Returns a `intervalId` for possible use with `clearInterval()`. Optionally, you can also pass arguments to the callback.

clearInterval(intervalId)
Stops a interval from triggering.

Child Processes

Node provides a tri-directional `popen(3)` facility through the `ChildProcess` class.

It is possible to stream data through the child's `stdin`, `stdout`, and `stderr` in a fully non-blocking way.

To create a child process use `require('child_process').spawn()`.

Child processes always have three streams associated with them. `child.stdin`, `child.stdout`, and `child.stderr`.

`ChildProcess` is an `EventEmitter`.

Event: 'exit'

`function (code, signal) {}`

This event is emitted after the child process ends. If the process terminated normally, `code` is the final exit code of the process, otherwise `null`. If the process terminated due to receipt of a signal, `signal` is the string name of the signal, otherwise `null`.

After this event is emitted, the `'output'` and `'error'` callbacks will no longer be made.

See `waitpid(2)`.

child.stdin

A `Writable Stream` that represents the child process's `stdin`. Closing this stream via `end()` often causes the child process to terminate.

child.stdout

A `Readable Stream` that represents the child process's `stdout`.

child.stderr

A `Readable Stream` that represents the child process's `stderr`.

child.pid

The PID of the child process.

Example:

```
var spawn = require('child_process').spawn,
    grep  = spawn('grep', ['ssh']);

console.log('Spawned child pid: ' + grep.pid);
grep.stdin.end();
```

child_process.spawn(command, args=[], [options])

Launches a new process with the given `command`, with command line arguments in `args`. If omitted, `args` defaults to an empty Array.

The third argument is used to specify additional options, which defaults to:

```
{ cwd: undefined
, env: process.env,
, customFds: [-1, -1, -1]
  }
```

`cwd` allows you to specify the working directory from which the process is spawned. Use `env` to specify environment variables that will be visible to the new process. With

`customFds` it is possible to hook up the new process' [stdin, stout, stderr] to existing streams; `-1` means that a new stream should be created.

Example of running `ls -lh /usr`, capturing `stdout`, `stderr`, and the exit code:

```
var sys   = require('sys'),
    spawn = require('child_process').spawn,
    ls    = spawn('ls', ['-lh', '/usr']);

ls.stdout.on('data', function (data) {
  sys.print('stdout: ' + data);
});

ls.stderr.on('data', function (data) {
  sys.print('stderr: ' + data);
});

ls.on('exit', function (code) {
  console.log('child process exited with code ' + code);
});
```

Example: A very elaborate way to run 'ps ax l grep ssh'

```
var sys   = require('sys'),
    spawn = require('child_process').spawn,
    ps    = spawn('ps', ['ax']),
    grep  = spawn('grep', ['ssh']);

ps.stdout.on('data', function (data) {
  grep.stdin.write(data);
});

ps.stderr.on('data', function (data) {
  sys.print('ps stderr: ' + data);
});

ps.on('exit', function (code) {
  if (code !== 0) {
    console.log('ps process exited with code ' + code);
  }
  grep.stdin.end();
});

grep.stdout.on('data', function (data) {
  sys.print(data);
});

grep.stderr.on('data', function (data) {
  sys.print('grep stderr: ' + data);
});
```

```
grep.on('exit', function (code) {
  if (code !== 0) {
    console.log('grep process exited with code ' + code);
  }
});
```

Example of checking for failed exec:

```
var spawn = require('child_process').spawn,
    child = spawn('bad_command');

child.stderr.on('data', function (data) {
  if (/^execvp\(\)/.test(data.asciiSlice(0,data.length))) {
    console.log('Failed to start child process.');
  }
});
```

See also: `child_process.exec()`

child_process.exec(command, [options], callback)
High-level way to execute a command as a child process, buffer the output, and return it all in a callback.

```
var sys   = require('sys'),
    exec  = require('child_process').exec,
    child;

child = exec('cat *.js bad_file | wc -l',
  function (error, stdout, stderr) {
    sys.print('stdout: ' + stdout);
    sys.print('stderr: ' + stderr);
    if (error !== null) {
      console.log('exec error: ' + error);
    }
});
```

The callback gets the arguments `(error, stdout, stderr)`. On success, `error` will be `null`. On error, `error` will be an instance of `Error` and `err.code` will be the exit code of the child process, and `err.signal` will be set to the signal that terminated the process.

There is a second optional argument to specify several options. The default options are

```
{ encoding: 'utf8'
, timeout: 0
, maxBuffer: 200*1024
, killSignal: 'SIGKILL'
, cwd: null
, env: null
}
```

If `timeout` is greater than 0, then it will kill the child process if it runs longer than `timeout` milliseconds. The child process is killed with `killsignal` (default: `'SIGKILL'`). `maxBuffer` specifies the largest amount of data allowed on stdout or stderr - if this value is exceeded then the child process is killed.

child.kill(signal='SIGTERM')

Send a signal to the child process. If no argument is given, the process will be sent `'SIGTERM'`. See `signal(7)` for a list of available signals.

```
var spawn = require('child_process').spawn,
    grep  = spawn('grep', ['ssh']);

grep.on('exit', function (code, signal) {
  console.log('child process terminated due to receipt of signal '+signal);
});

// send SIGHUP to process
grep.kill('SIGHUP');
```

Note that while the function is called `kill`, the signal delivered to the child process may not actually kill it. `kill` really just sends a signal to a process.

See `kill(2)`

Script

`Script` class compiles and runs JavaScript code. You can access this class with:

```
var Script = process.binding('evals').Script;
```

New JavaScript code can be compiled and run immediately or compiled, saved, and run later.

Script.runInThisContext(code, [filename])

Similar to `process.compile`. `Script.runInThisContext` compiles `code` as if it were loaded from `filename`, runs it and returns the result. Running code does not have access to local scope. `filename` is optional.

Example of using `Script.runInThisContext` and `eval` to run the same code:

```
var localVar = 123,
    usingscript, evaled,
    Script = process.binding('evals').Script;

usingscript = Script.runInThisContext('localVar = 1;',
  'myfile.js');
console.log('localVar: ' + localVar + ', usingscript: ' +
```

```
    usingscript);
evaled = eval('localVar = 1;');
console.log('localVar: ' + localVar + ', evaled: ' +
   evaled);

// localVar: 123, usingscript: 1
// localVar: 1, evaled: 1
```

`Script.runInThisContext` does not have access to the local scope, so `localVar` is unchanged. `eval` does have access to the local scope, so `localVar` is changed.

In case of syntax error in `code`, `Script.runInThisContext` emits the syntax error to stderr and throws.an exception.

Script.runInNewContext(code, [sandbox], [filename])

`Script.runInNewContext` compiles `code` to run in `sandbox` as if it were loaded from `filename`, then runs it and returns the result. Running code does not have access to local scope and the object `sandbox` will be used as the global object for `code`. `sandbox` and `filename` are optional.

Example: compile and execute code that increments a global variable and sets a new one. These globals are contained in the sandbox.

```
var sys = require('sys'),
    Script = process.binding('evals').Script,
    sandbox = {
      animal: 'cat',
      count: 2
    };

Script.runInNewContext(
  'count += 1; name = "kitty"', sandbox, 'myfile.js');
console.log(sys.inspect(sandbox));

// { animal: 'cat', count: 3, name: 'kitty' }
```

Note that running untrusted code is a tricky business requiring great care. To prevent accidental global variable leakage, `Script.runInNewContext` is quite useful, but safely running untrusted code requires a separate process.

In case of syntax error in `code`, `Script.runInThisContext` emits the syntax error to stderr and throws an exception.

new Script(code, [filename])

`new Script` compiles `code` as if it were loaded from `filename`, but does not run it. Instead, it returns a `script` object representing this compiled code. This script can be run later many times using methods below. The returned script is not bound to any global object. It is

bound before each run, just for that run. `filename` is optional.

In case of syntax error in `code`, `new Script` emits the syntax error to stderr and throws an exception.

script.runInThisContext()

Similar to `Script.runInThisContext` (note capital 'S'), but now being a method of a precompiled Script object. `script.runInThisContext` runs the code of `script` and returns the result. Running code does not have access to local scope, but does have access to the `global` object (v8: in actual context).

Example of using `script.runInThisContext` to compile code once and run it multiple times:

```
var Script = process.binding('evals').Script,
    scriptObj, i;

globalVar = 0;

scriptObj = new Script('globalVar += 1', 'myfile.js');

for (i = 0; i < 1000 ; i += 1) {
  scriptObj.runInThisContext();
}

console.log(globalVar);

// 1000
```

script.runInNewContext([sandbox])

Similar to `Script.runInNewContext` (note capital 'S'), but now being a method of a precompiled Script object. `script.runInNewContext` runs the code of `script` with `sandbox` as the global object and returns the result. Running code does not have access to local scope. `sandbox` is optional.

Example: compile code that increments a global variable and sets one, then execute this code multiple times. These globals are contained in the sandbox.

```
var sys = require('sys'),
    Script = process.binding('evals').Script,
    scriptObj, i,
    sandbox = {
      animal: 'cat',
      count: 2
    };

scriptObj = new Script(
    'count += 1; name = "kitty"', 'myfile.js');
```

```
  for (i = 0; i < 10 ; i += 1) {
    scriptObj.runInNewContext(sandbox);
  }

  console.log(sys.inspect(sandbox));

  // { animal: 'cat', count: 12, name: 'kitty' }
```

Note that running untrusted code is a tricky business requiring great care. To prevent accidental global variable leakage, `script.runInNewContext` is quite useful, but safely running untrusted code requires a separate process.

File System

File I/O is provided by simple wrappers around standard POSIX functions. To use this module do `require('fs')`. All the methods have asynchronous and synchronous forms.

The asynchronous form always take a completion callback as its last argument. The arguments passed to the completion callback depend on the method, but the first argument is always reserved for an exception. If the operation was completed successfully, then the first argument will be `null` or `undefined`.

Here is an example of the asynchronous version:

```
var fs = require('fs');

fs.unlink('/tmp/hello', function (err) {
  if (err) throw err;
  console.log('successfully deleted /tmp/hello');
});
```

Here is the synchronous version:

```
var fs = require('fs');

fs.unlinkSync('/tmp/hello')
console.log('successfully deleted /tmp/hello');
```

With the asynchronous methods there is no guaranteed ordering. So the following is prone to error:

```
fs.rename('/tmp/hello', '/tmp/world', function (err) {
  if (err) throw err;
  console.log('renamed complete');
});
fs.stat('/tmp/world', function (err, stats) {
  if (err) throw err;
  console.log('stats: ' + JSON.stringify(stats));
```

```
  });
```

It could be that `fs.stat` is executed before `fs.rename`. The correct way to do this is to chain the callbacks.

```
fs.rename('/tmp/hello', '/tmp/world', function (err) {
  if (err) throw err;
  fs.stat('/tmp/world', function (err, stats) {
    if (err) throw err;
    console.log('stats: ' + JSON.stringify(stats));
  });
});
```

In busy processes, the programmer is <u>strongly encouraged</u> to use the asynchronous versions of these calls. The synchronous versions will block the entire process until they complete--halting all connections.

fs.rename(path1, path2, [callback])
Asynchronous rename(2). No arguments other than a possible exception are given to the completion callback.

fs.renameSync(path1, path2)
Synchronous rename(2).

fs.truncate(fd, len, [callback])
Asynchronous ftruncate(2). No arguments other than a possible exception are given to the completion callback.

fs.truncateSync(fd, len)
Synchronous ftruncate(2).

fs.chmod(path, mode, [callback])
Asynchronous chmod(2). No arguments other than a possible exception are given to the completion callback.

fs.chmodSync(path, mode)
Synchronous chmod(2).

fs.stat(path, [callback])
Asynchronous stat(2). The callback gets two arguments `(err, stats)` where `stats` is a `fs.Stats` object. It looks like this:

```
{ dev: 2049
, ino: 305352
, mode: 16877
, nlink: 12
, uid: 1000
```

```
, gid: 1000
, rdev: 0
, size: 4096
, blksize: 4096
, blocks: 8
, atime: '2009-06-29T11:11:55Z'
, mtime: '2009-06-29T11:11:40Z'
, ctime: '2009-06-29T11:11:40Z'
}
```

See the `fs.Stats` section below for more information.

fs.lstat(path, [callback])

Asynchronous lstat(2). The callback gets two arguments `(err, stats)` where `stats` is a `fs.Stats` object.

fs.fstat(fd, [callback])

Asynchronous fstat(2). The callback gets two arguments `(err, stats)` where `stats` is a `fs.Stats` object.

fs.statSync(path)

Synchronous stat(2). Returns an instance of `fs.Stats`.

fs.lstatSync(path)

Synchronous lstat(2). Returns an instance of `fs.Stats`.

fs.fstatSync(fd)

Synchronous fstat(2). Returns an instance of `fs.Stats`.

fs.link(srcpath, dstpath, [callback])

Asynchronous link(2). No arguments other than a possible exception are given to the completion callback.

fs.linkSync(dstpath, srcpath)

Synchronous link(2).

fs.symlink(linkdata, path, [callback])

Asynchronous symlink(2). No arguments other than a possible exception are given to the completion callback.

fs.symlinkSync(linkdata, path)

Synchronous symlink(2).

fs.readlink(path, [callback])

Asynchronous readlink(2). The callback gets two arguments `(err, resolvedPath)`.

fs.readlinkSync(path)

Synchronous readlink(2). Returns the resolved path.

fs.realpath(path, [callback])
Asynchronous realpath(2). The callback gets two arguments `(err, resolvedPath)`.

fs.realpathSync(path)
Synchronous realpath(2). Returns the resolved path.

fs.unlink(path, [callback])
Asynchronous unlink(2). No arguments other than a possible exception are given to the completion callback.

fs.unlinkSync(path)
Synchronous unlink(2).

fs.rmdir(path, [callback])
Asynchronous rmdir(2). No arguments other than a possible exception are given to the completion callback.

fs.rmdirSync(path)
Synchronous rmdir(2).

fs.mkdir(path, mode, [callback])
Asynchronous mkdir(2). No arguments other than a possible exception are given to the completion callback.

fs.mkdirSync(path, mode)
Synchronous mkdir(2).

fs.readdir(path, [callback])
Asynchronous readdir(3). Reads the contents of a directory. The callback gets two arguments `(err, files)` where `files` is an array of the names of the files in the directory excluding `'.'` and `'..'`.

fs.readdirSync(path)
Synchronous readdir(3). Returns an array of filenames excluding `'.'` and `'..'`.

fs.close(fd, [callback])
Asynchronous close(2). No arguments other than a possible exception are given to the completion callback.

fs.closeSync(fd)
Synchronous close(2).

fs.open(path, flags, mode=0666, [callback])
Asynchronous file open. See open(2). Flags can be 'r', 'r+', 'w', 'w+', 'a', or 'a+'. The callback

gets two arguments `(err, fd)`.

fs.openSync(path, flags, mode=0666)
Synchronous open(2).

fs.write(fd, buffer, offset, length, position, [callback])
Write `buffer` to the file specified by `fd`.

`offset` and `length` determine the part of the buffer to be written.

`position` refers to the offset from the beginning of the file where this data should be written. If `position` is `null`, the data will be written at the current position. See pwrite(2).

The callback will be given two arguments `(err, written)` where `written` specifies how many bytes were written.

fs.writeSync(fd, buffer, offset, length, position)
Synchronous version of buffer-based `fs.write()`. Returns the number of bytes written.

fs.writeSync(fd, str, position, encoding='utf8')
Synchronous version of string-based `fs.write()`. Returns the number of bytes written.

fs.read(fd, buffer, offset, length, position, [callback])
Read data from the file specified by `fd`.

`buffer` is the buffer that the data will be written to.

`offset` is offset within the buffer where writing will start.

`length` is an integer specifying the number of bytes to read.

`position` is an integer specifying where to begin reading from in the file. If `position` is `null`, data will be read from the current file position.

The callback is given the two arguments, `(err, bytesRead)`.

fs.readSync(fd, buffer, offset, length, position)
Synchronous version of buffer-based `fs.read`. Returns the number of `bytesRead`.

fs.readSync(fd, length, position, encoding)
Synchronous version of string-based `fs.read`. Returns the number of `bytesRead`.

fs.readFile(filename, [encoding], [callback])
Asynchronously reads the entire contents of a file. Example:

```
fs.readFile('/etc/passwd', function (err, data) {
  if (err) throw err;
  console.log(data);
```

```
    });
```

The callback is passed two arguments `(err, data)`, where `data` is the contents of the file.

If no encoding is specified, then the raw buffer is returned.

fs.readFileSync(filename, [encoding])

Synchronous version of `fs.readFile`. Returns the contents of the `filename`.

If `encoding` is specified then this function returns a string. Otherwise it returns a buffer.

fs.writeFile(filename, data, encoding='utf8', [callback])

Asynchronously writes data to a file. `data` can be a string or a buffer.

Example:

```
fs.writeFile('message.txt', 'Hello Node', function (err) {
  if (err) throw err;
  console.log('It\'s saved!');
});
```

fs.writeFileSync(filename, data, encoding='utf8')

The synchronous version of `fs.writeFile`.

fs.watchFile(filename, [options], listener)

Watch for changes on `filename`. The callback `listener` will be called each time the file changes.

The second argument is optional. The `options` if provided should be an object containing two members a boolean, `persistent`, and `interval`, a polling value in milliseconds. The default is `{persistent: true, interval: 0}`.

The `listener` gets two arguments the current stat object and the previous stat object:

```
fs.watchFile(f, function (curr, prev) {
  console.log('the current mtime is: ' + curr.mtime);
  console.log('the previous mtime was: ' + prev.mtime);
});
```

These stat objects are instances of `fs.Stat`.

fs.unwatchFile(filename)

Stop watching for changes on `filename`.

fs.Stats

Objects returned from `fs.stat()` and `fs.lstat()` are of this type.

- `stats.isFile()`
- `stats.isDirectory()`
- `stats.isBlockDevice()`
- `stats.isCharacterDevice()`
- `stats.isSymbolicLink()` (only valid with `fs.lstat()`)
- `stats.isFIFO()`
- `stats.isSocket()`

fs.ReadStream

`ReadStream` is a `Readable Stream`.

fs.createReadStream(path, [options])

Returns a new ReadStream object (See `Readable Stream`).

`options` is an object with the following defaults:

```
{ 'flags': 'r'
, 'encoding': null
, 'mode': 0666
, 'bufferSize': 4 * 1024
}
```

`options` can include `start` and `end` values to read a range of bytes from the file instead of the entire file. Both `start` and `end` are inclusive and start at 0. When used, both the limits must be specified always.

An example to read the last 10 bytes of a file which is 100 bytes long:

```
fs.createReadStream('sample.txt', {start: 90, end: 99});
```

fs.WriteStream

`WriteStream` is a `Writable Stream`.

Event: 'open'
`function (fd) { }`

`fd` is the file descriptor used by the WriteStream.

fs.createWriteStream(path, [options])

Returns a new WriteStream object (See `Writable Stream`).

`options` is an object with the following defaults:

```
{ 'flags': 'w'
```

```
, 'encoding': null
, 'mode': 0666
}
```

HTTP

To use the HTTP server and client one must `require('http')`.

The HTTP interfaces in Node are designed to support many features of the protocol which have been traditionally difficult to use. In particular, large, possibly chunk-encoded, messages. The interface is careful to never buffer entire requests or responses--the user is able to stream data.

HTTP message headers are represented by an object like this:

```
{ 'content-length': '123'
, 'content-type': 'text/plain'
, 'stream': 'keep-alive'
, 'accept': '*/*'
}
```

Keys are lowercased. Values are not modified.

In order to support the full spectrum of possible HTTP applications, Node's HTTP API is very low-level. It deals with stream handling and message parsing only. It parses a message into headers and body but it does not parse the actual headers or the body.

HTTPS is supported if OpenSSL is available on the underlying platform.

http.Server

This is an `EventEmitter` with the following events:

Event: 'request'
`function (request, response) { }`

`request` is an instance of `http.ServerRequest` and `response` is an instance of `http.ServerResponse`

Event: 'connection'
`function (stream) { }`

When a new TCP stream is established. `stream` is an object of type `net.Stream`. Usually users will not want to access this event. The `stream` can also be accessed at `request.connection`.

Event: 'close'
`function (errno) { }`

Emitted when the server closes.

Event: 'request'

```
function (request, response) {}
```

Emitted each time there is request. Note that there may be multiple requests per connection (in the case of keep-alive connections).

Event: 'upgrade'

```
function (request, socket, head)
```

Emitted each time a client requests a http upgrade. If this event isn't listened for, then clients requesting an upgrade will have their connections closed.

- `request` is the arguments for the http request, as it is in the request event.
- `socket` is the network socket between the server and client.
- `head` is an instance of Buffer, the first packet of the upgraded stream, this may be empty.

After this event is emitted, the request's socket will not have a `data` event listener, meaning you will need to bind to it in order to handle data sent to the server on that socket.

Event: 'clientError'

```
function (exception) {}
```

If a client connection emits an 'error' event - it will forwarded here.

http.createServer(requestListener)

Returns a new web server object.

The `requestListener` is a function which is automatically added to the `'request'` event.

server.listen(port, [hostname], [callback])

Begin accepting connections on the specified port and hostname. If the hostname is omitted, the server will accept connections directed to any IPv4 address (`INADDR_ANY`).

To listen to a unix socket, supply a filename instead of port and hostname.

This function is asynchronous. The last parameter `callback` will be called when the server has been bound to the port.

server.listen(path, [callback])

Start a UNIX socket server listening for connections on the given `path`.

This function is asynchronous. The last parameter `callback` will be called when the server has been bound.

server.setSecure(credentials)

Enables HTTPS support for the server, with the crypto module credentials specifying the private key and certificate of the server, and optionally the CA certificates for use in client authentication.

If the credentials hold one or more CA certificates, then the server will request for the client to submit a client certificate as part of the HTTPS connection handshake. The validity and content of this can be accessed via verifyPeer() and getPeerCertificate() from the server's request.connection.

server.close()

Stops the server from accepting new connections.

http.ServerRequest

This object is created internally by a HTTP server--not by the user--and passed as the first argument to a `'request'` listener.

This is an `EventEmitter` with the following events:

Event: 'data'

```
function (chunk) { }
```

Emitted when a piece of the message body is received.

Example: A chunk of the body is given as the single argument. The transfer-encoding has been decoded. The body chunk is a string. The body encoding is set with `request.setBodyEncoding()`.

Event: 'end'

```
function () { }
```

Emitted exactly once for each message. No arguments. After emitted no other events will be emitted on the request.

request.method

The request method as a string. Read only. Example: `'GET'`, `'DELETE'`.

request.url

Request URL string. This contains only the URL that is present in the actual HTTP request. If the request is:

```
GET /status?name=ryan HTTP/1.1\r\n
Accept: text/plain\r\n
\r\n
```

Then `request.url` will be:

```
'/status?name=ryan'
```

If you would like to parse the URL into its parts, you can use
`require('url').parse(request.url)`. Example:

```
node> require('url').parse('/status?name=ryan')
{ href: '/status?name=ryan'
, search: '?name=ryan'
, query: 'name=ryan'
, pathname: '/status'
}
```

If you would like to extract the params from the query string, you can use the
`require('querystring').parse` function, or pass `true` as the second argument to
`require('url').parse`. Example:

```
node> require('url').parse('/status?name=ryan', true)
{ href: '/status?name=ryan'
, search: '?name=ryan'
, query: { name: 'ryan' }
, pathname: '/status'
}
```

request.headers
Read only.

request.httpVersion
The HTTP protocol version as a string. Read only. Examples: `'1.1'`, `'1.0'`. Also
`request.httpVersionMajor` is the first integer and `request.httpVersionMinor` is the
second.

request.setEncoding(encoding=null)
Set the encoding for the request body. Either `'utf8'` or `'binary'`. Defaults to `null`, which
means that the `'data'` event will emit a `Buffer` object..

request.pause()
Pauses request from emitting events. Useful to throttle back an upload.

request.resume()
Resumes a paused request.

request.connection
The `net.Stream` object associated with the connection.

With HTTPS support, use request.connection.verifyPeer() and

request.connection.getPeerCertificate() to obtain the client's authentication details.

http.ServerResponse

This object is created internally by a HTTP server--not by the user. It is passed as the second parameter to the `'request'` event. It is a `Writable Stream`.

response.writeHead(statusCode, [reasonPhrase], [headers])
Sends a response header to the request. The status code is a 3-digit HTTP status code, like `404`. The last argument, `headers`, are the response headers. Optionally one can give a human-readable `reasonPhrase` as the second argument.

Example:

```
var body = 'hello world';
response.writeHead(200, {
  'Content-Length': body.length,
  'Content-Type': 'text/plain'
});
```

This method must only be called once on a message and it must be called before `response.end()` is called.

response.write(chunk, encoding='utf8')
This method must be called after `writeHead` was called. It sends a chunk of the response body. This method may be called multiple times to provide successive parts of the body.

`chunk` can be a string or a buffer. If `chunk` is a string, the second parameter specifies how to encode it into a byte stream. By default the `encoding` is `'utf8'`.

Note: This is the raw HTTP body and has nothing to do with higher-level multi-part body encodings that may be used.

The first time `response.write()` is called, it will send the buffered header information and the first body to the client. The second time `response.write()` is called, Node assumes you're going to be streaming data, and sends that separately. That is, the response is buffered up to the first chunk of body.

response.addTrailers(headers)
This method adds HTTP trailing headers (a header but at the end of the message) to the response.

Trailers will **only** be emitted if chunked encoding is used for the response; if it is not (e.g., if the request was HTTP/1.0), they will be silently discarded.

Note that HTTP requires the `Trailer` header to be sent if you intend to emit trailers, with a

list of the header fields in its value. E.g.,

```
response.writeHead(200, { 'Content-Type': 'text/plain',
                          'Trailer': 'TraceInfo' });
response.write(fileData);
response.addTrailers({'Content-MD5': "7895bf4b8828b55ceaf47747b4bca667"});
response.end();
```

response.end([data], [encoding])

This method signals to the server that all of the response headers and body has been sent; that server should consider this message complete. The method, `response.end()`, MUST be called on each response.

If `data` is specified, it is equivalent to calling `response.write(data, encoding)` followed by `response.end()`.

http.Client

An HTTP client is constructed with a server address as its argument, the returned handle is then used to issue one or more requests. Depending on the server connected to, the client might pipeline the requests or reestablish the stream after each stream. Currently the implementation does not pipeline requests.

Example of connecting to `google.com`:

```
var http = require('http');
var google = http.createClient(80, 'www.google.com');
var request = google.request('GET', '/',
  {'host': 'www.google.com'});
request.end();
request.on('response', function (response) {
  console.log('STATUS: ' + response.statusCode);
  console.log('HEADERS: ' + JSON.stringify(response.headers));
  response.setEncoding('utf8');
  response.on('data', function (chunk) {
    console.log('BODY: ' + chunk);
  });
});
```

There are a few special headers that should be noted.

- The 'Host' header is not added by Node, and is usually required by website.

- Sending a 'Connection: keep-alive' will notify Node that the connection to the server should be persisted until the next request.

- Sending a 'Content-length' header will disable the default chunked encoding.

Event: 'upgrade'
```
function (request, socket, head)
```

Emitted each time a server responds to a request with an upgrade. If this event isn't being listened for, clients receiving an upgrade header will have their connections closed.

See the description of the `upgrade` event for `http.Server` for further details.

http.createClient(port, host='localhost', secure=false, [credentials])
Constructs a new HTTP client. `port` and `host` refer to the server to be connected to. A stream is not established until a request is issued.

`secure` is an optional boolean flag to enable https support and `credentials` is an optional credentials object from the crypto module, which may hold the client's private key, certificate, and a list of trusted CA certificates.

If the connection is secure, but no explicit CA certificates are passed in the credentials, then node.js will default to the publicly trusted list of CA certificates, as given in http://mxr.mozilla.org/mozilla/source/security/nss/lib/ckfw/builtins/certdata.txt

client.request(method='GET', path, [request_headers])
Issues a request; if necessary establishes stream. Returns a `http.ClientRequest` instance.

`method` is optional and defaults to 'GET' if omitted.

`request_headers` is optional. Additional request headers might be added internally by Node. Returns a `ClientRequest` object.

Do remember to include the `Content-Length` header if you plan on sending a body. If you plan on streaming the body, perhaps set `Transfer-Encoding: chunked`.

NOTE: the request is not complete. This method only sends the header of the request. One needs to call `request.end()` to finalize the request and retrieve the response. (This sounds convoluted but it provides a chance for the user to stream a body to the server with `request.write()`.)

client.verifyPeer()
Returns true or false depending on the validity of the server's certificate in the context of the defined or default list of trusted CA certificates.

client.getPeerCertificate()
Returns a JSON structure detailing the server's certificate, containing a dictionary with keys for the certificate 'subject', 'issuer', 'valid_from' and 'valid_to'

http.ClientRequest

This object is created internally and returned from the `request()` method of a `http.Client`. It represents an in-progress request whose header has already been sent.

To get the response, add a listener for `'response'` to the request object. `'response'` will be emitted from the request object when the response headers have been received. The `'response'` event is executed with one argument which is an instance of `http.ClientResponse`.

During the `'response'` event, one can add listeners to the response object; particularly to listen for the `'data'` event. Note that the `'response'` event is called before any part of the response body is received, so there is no need to worry about racing to catch the first part of the body. As long as a listener for `'data'` is added during the `'response'` event, the entire body will be caught.

```
// Good
request.on('response', function (response) {
  response.on('data', function (chunk) {
    console.log('BODY: ' + chunk);
  });
});

// Bad - misses all or part of the body
request.on('response', function (response) {
  setTimeout(function () {
    response.on('data', function (chunk) {
      console.log('BODY: ' + chunk);
    });
  }, 10);
});
```

This is a `Writable Stream`.

This is an `EventEmitter` with the following events:

Event 'response'
`function (response) { }`

Emitted when a response is received to this request. This event is emitted only once. The `response` argument will be an instance of `http.ClientResponse`.

request.write(chunk, encoding='utf8')
Sends a chunk of the body. By calling this method many times, the user can stream a request body to a server--in that case it is suggested to use the `['Transfer-Encoding', 'chunked']` header line when creating the request.

The `chunk` argument should be an array of integers or a string.

The `encoding` argument is optional and only applies when `chunk` is a string.

request.end([data], [encoding])

Finishes sending the request. If any parts of the body are unsent, it will flush them to the stream. If the request is chunked, this will send the terminating `'0\r\n\r\n'`.

If `data` is specified, it is equivalent to calling `request.write(data, encoding)` followed by `request.end()`.

http.ClientResponse

This object is created when making a request with `http.Client`. It is passed to the `'response'` event of the request object.

The response implements the `Readable Stream` interface.

Event: 'data'

```
function (chunk) {}
```

Emitted when a piece of the message body is received.

```
Example: A chunk of the body is given as the single
argument. The transfer-encoding has been decoded.  The
body chunk a String.  The body encoding is set with
 `response.setBodyEncoding()`.
```

Event: 'end'

```
function () {}
```

Emitted exactly once for each message. No arguments. After emitted no other events will be emitted on the response.

response.statusCode

The 3-digit HTTP response status code. E.G. `404`.

response.httpVersion

The HTTP version of the connected-to server. Probably either `'1.1'` or `'1.0'`. Also `response.httpVersionMajor` is the first integer and `response.httpVersionMinor` is the second.

response.headers

The response headers object.

response.setEncoding(encoding=null)

Set the encoding for the response body. Either `'utf8'`, `'ascii'`, or `'base64'`. Defaults to `null`, which means that the `'data'` event will emit a `Buffer` object..

response.pause()

Pauses response from emitting events. Useful to throttle back a download.

response.resume()

Resumes a paused response.

response.client

A reference to the `http.Client` that this response belongs to.

net.Server

This class is used to create a TCP or UNIX server.

Here is an example of a echo server which listens for connections on port 8124:

```
var net = require('net');
var server = net.createServer(function (stream) {
  stream.setEncoding('utf8');
  stream.on('connect', function () {
    stream.write('hello\r\n');
  });
  stream.on('data', function (data) {
    stream.write(data);
  });
  stream.on('end', function () {
    stream.write('goodbye\r\n');
    stream.end();
  });
});
server.listen(8124, 'localhost');
```

To listen on the socket `'/tmp/echo.sock'`, the last line would just be changed to

```
server.listen('/tmp/echo.sock');
```

This is an `EventEmitter` with the following events:

Event: 'connection'

`function (stream) {}`

Emitted when a new connection is made. `stream` is an instance of `net.Stream`.

Event: 'close'

`function () {}`

Emitted when the server closes.

net.createServer(connectionListener)

Creates a new TCP server. The `connectionListener` argument is automatically set as a

listener for the `'connection'` event.

server.listen(port, [host], [callback])

Begin accepting connections on the specified `port` and `host`. If the `host` is omitted, the server will accept connections directed to any IPv4 address (`INADDR_ANY`).

This function is asynchronous. The last parameter `callback` will be called when the server has been bound.

server.listen(path, [callback])

Start a UNIX socket server listening for connections on the given `path`.

This function is asynchronous. The last parameter `callback` will be called when the server has been bound.

server.listenFD(fd)

Start a server listening for connections on the given file descriptor.

This file descriptor must have already had the `bind(2)` and `listen(2)` system calls invoked on it.

server.close()

Stops the server from accepting new connections. This function is asynchronous, the server is finally closed when the server emits a `'close'` event.

server.maxConnections

Set this property to reject connections when the server's connection count gets high.

server.connections

The number of concurrent connections on the server.

net.isIP

net.isIP(input)

Tests if input is an IP address. Returns 0 for invalid strings, returns 4 for IP version 4 addresses, and returns 6 for IP version 6 addresses.

net.isIPv4(input)

Returns true if input is a version 4 IP address, otherwise returns false.

net.isIPv6(input)

Returns true if input is a version 6 IP address, otherwise returns false.

net.Stream

This object is an abstraction of of a TCP or UNIX socket. `net.stream` instance implement a

duplex stream interface. They can be created by the user and used as a client (with `connect()`) or they can be created by Node and passed to the user through the `'connection'` event of a server.

`net.Stream` instances are EventEmitters with the following events:

Event: 'connect'

```
function () { }
```

Emitted when a stream connection successfully is established. See `connect()`.

Event: 'secure'

```
function () { }
```

Emitted when a stream connection successfully establishes an SSL handshake with its peer.

Event: 'data'

```
function (data) { }
```

Emitted when data is received. The argument `data` will be a `Buffer` or `String`. Encoding of data is set by `stream.setEncoding()`. (See the section on `Readable Stream` for more information.)

Event: 'end'

```
function () { }
```

Emitted when the other end of the stream sends a FIN packet. After this is emitted the `readyState` will be `'writeOnly'`. One should probably just call `stream.end()` when this event is emitted.

Event: 'timeout'

```
function () { }
```

Emitted if the stream times out from inactivity. This is only to notify that the stream has been idle. The user must manually close the connection.

See also: `stream.setTimeout()`

Event: 'drain'

```
function () { }
```

Emitted when the write buffer becomes empty. Can be used to throttle uploads.

Event: 'error'

```
function (exception) { }
```

Emitted when an error occurs. The `'close'` event will be called directly following this

Event: 'close'
```
function (had_error) { }
```

Emitted once the stream is fully closed. The argument `had_error` is a boolean which says if the stream was closed due to a transmission error.

net.createConnection(port, host='127.0.0.1')
Construct a new stream object and opens a stream to the specified `port` and `host`. If the second parameter is omitted, localhost is assumed.

When the stream is established the `'connect'` event will be emitted.

stream.connect(port, host='127.0.0.1')
Opens a stream to the specified `port` and `host`. `createConnection()` also opens a stream; normally this method is not needed. Use this only if a stream is closed and you want to reuse the object to connect to another server.

This function is asynchronous. When the `'connect'` event is emitted the stream is established. If there is a problem connecting, the `'connect'` event will not be emitted, the `'error'` event will be emitted with the exception.

stream.remoteAddress
The string representation of the remote IP address. For example, `'74.125.127.100'` or `'2001:4860:a005::68'`.

This member is only present in server-side connections.

stream.readyState
Either `'closed'`, `'open'`, `'opening'`, `'readOnly'`, or `'writeOnly'`.

stream.setEncoding(encoding=null)
Sets the encoding (either `'ascii'`, `'utf8'`, or `'base64'`) for data that is received.

stream.setSecure([credentials])
Enables SSL support for the stream, with the crypto module credentials specifying the private key and certificate of the stream, and optionally the CA certificates for use in peer authentication.

If the credentials hold one ore more CA certificates, then the stream will request for the peer to submit a client certificate as part of the SSL connection handshake. The validity and content of this can be accessed via verifyPeer() and getPeerCertificate().

stream.verifyPeer()
Returns true or false depending on the validity of the peers's certificate in the context of

the defined or default list of trusted CA certificates.

stream.getPeerCertificate()

Returns a JSON structure detailing the peer's certificate, containing a dictionary with keys for the certificate 'subject', 'issuer', 'valid_from' and 'valid_to'

stream.write(data, encoding='ascii')

Sends data on the stream. The second parameter specifies the encoding in the case of a string--it defaults to ASCII because encoding to UTF8 is rather slow.

Returns `true` if the entire data was flushed successfully to the kernel buffer. Returns `false` if all or part of the data was queued in user memory. `'drain'` will be emitted when the buffer is again free.

stream.end([data], [encoding])

Half-closes the stream. I.E., it sends a FIN packet. It is possible the server will still send some data. After calling this `readyState` will be `'readOnly'`.

If `data` is specified, it is equivalent to calling `stream.write(data, encoding)` followed by `stream.end()`.

stream.destroy()

Ensures that no more I/O activity happens on this stream. Only necessary in case of errors (parse error or so).

stream.pause()

Pauses the reading of data. That is, `'data'` events will not be emitted. Useful to throttle back an upload.

stream.resume()

Resumes reading after a call to `pause()`.

stream.setTimeout(timeout)

Sets the stream to timeout after `timeout` milliseconds of inactivity on the stream. By default `net.stream` do not have a timeout.

When an idle timeout is triggered the stream will receive a `'timeout'` event but the connection will not be severed. The user must manually `end()` or `destroy()` the stream.

If `timeout` is 0, then the existing idle timeout is disabled.

stream.setNoDelay(noDelay=true)

Disables the Nagle algorithm. By default TCP connections use the Nagle algorithm, they buffer data before sending it off. Setting `noDelay` will immediately fire off data each time `stream.write()` is called.

stream.setKeepAlive(enable=false, [initialDelay])

Enable/disable keep-alive functionality, and optionally set the initial delay before the first keepalive probe is sent on an idle stream. Set `initialDelay` (in milliseconds) to set the delay between the last data packet received and the first keepalive probe. Setting 0 for initialDelay will leave the value unchanged from the default (or previous) setting.

Crypto

Use `require('crypto')` to access this module.

The crypto module requires OpenSSL to be available on the underlying platform. It offers a way of encapsulating secure credentials to be used as part of a secure HTTPS net or http connection.

It also offers a set of wrappers for OpenSSL's hash, hmac, cipher, decipher, sign and verify methods.

crypto.createCredentials(details)

Creates a credentials object, with the optional details being a dictionary with keys:

`key` : a string holding the PEM encoded private key

`cert` : a string holding the PEM encoded certificate

`ca` : either a string or list of strings of PEM encoded CA certificates to trust.

If no 'ca' details are given, then node.js will use the default publicly trusted list of CAs as given in http://mxr.mozilla.org/mozilla/source/security/nss/lib/ckfw/builtins/certdata.txt

crypto.createHash(algorithm)

Creates and returns a hash object, a cryptographic hash with the given algorithm which can be used to generate hash digests.

`algorithm` is dependent on the available algorithms supported by the version of OpenSSL on the platform. Examples are sha1, md5, sha256, sha512, etc. On recent releases, `openssl list-message-digest-algorithms` will display the available digest algorithms.

hash.update(data)

Updates the hash content with the given `data`. This can be called many times with new data as it is streamed.

hash.digest(encoding='binary')

Calculates the digest of all of the passed data to be hashed. The `encoding` can be 'hex', 'binary' or 'base64'.

crypto.createHmac(algorithm, key)

Creates and returns a hmac object, a cryptographic hmac with the given algorithm and key.

`algorithm` is dependent on the available algorithms supported by OpenSSL - see createHash above. `key` is the hmac key to be used.

hmac.update(data)

Update the hmac content with the given `data`. This can be called many times with new data as it is streamed.

hmac.digest(encoding='binary')

Calculates the digest of all of the passed data to the hmac. The `encoding` can be 'hex', 'binary' or 'base64'.

crypto.createCipher(algorithm, key)

Creates and returns a cipher object, with the given algorithm and key.

`algorithm` is dependent on OpenSSL, examples are aes192, etc. On recent releases, `openssl list-cipher-algorithms` will display the available cipher algorithms.

cipher.update(data, input_encoding='binary', output_encoding='binary')

Updates the cipher with `data`, the encoding of which is given in `input_encoding` and can be 'utf8', 'ascii' or 'binary'. The `output_encoding` specifies the output format of the enciphered data, and can be 'binary', 'base64' or 'hex'.

Returns the enciphered contents, and can be called many times with new data as it is streamed.

cipher.final(output_encoding='binary')

Returns any remaining enciphered contents, with `output_encoding` being one of: 'binary', 'ascii' or 'utf8'.

crypto.createDecipher(algorithm, key)

Creates and returns a decipher object, with the given algorithm and key. This is the mirror of the cipher object above.

decipher.update(data, input_encoding='binary', output_encoding='binary')

Updates the decipher with `data`, which is encoded in 'binary', 'base64' or 'hex'. The `output_decoding` specifies in what format to return the deciphered plaintext - either 'binary', 'ascii' or 'utf8'.

decipher.final(output_encoding='binary')

Returns any remaining plaintext which is deciphered, with `output_encoding' being one of: 'binary', 'ascii' or 'utf8'.

crypto.createSign(algorithm)

Creates and returns a signing object, with the given algorithm. On recent OpenSSL releases, `openssl list-public-key-algorithms` will display the available signing algorithms. Examples are 'RSA-SHA256'.

signer.update(data)

Updates the signer object with data. This can be called many times with new data as it is streamed.

signer.sign(private_key, output_format='binary')

Calculates the signature on all the updated data passed through the signer. `private_key` is a string containing the PEM encoded private key for signing.

Returns the signature in `output_format` which can be 'binary', 'hex' or 'base64'

crypto.createVerify(algorithm)

Creates and returns a verification object, with the given algorithm. This is the mirror of the signing object above.

verifier.update(data)

Updates the verifyer object with data. This can be called many times with new data as it is streamed.

verifier.verify(public_key, signature, signature_format='binary')

Verifies the signed data by using the `public_key` which is a string containing the PEM encoded public key, and `signature`, which is the previously calculates signature for the data, in the `signature_format` which can be 'binary', 'hex' or 'base64'.

Returns true or false depending on the validity of the signature for the data and public key.

DNS

Use `require('dns')` to access this module.

Here is an example which resolves `'www.google.com'` then reverse resolves the IP addresses which are returned.

```
var dns = require('dns');

dns.resolve4('www.google.com', function (err, addresses) {
  if (err) throw err;

  console.log('addresses: ' + JSON.stringify(addresses));

  addresses.forEach(function (a) {
    dns.reverse(a, function (err, domains) {
      if (err) {
```

```
            console.log('reverse for ' + a + ' failed: ' +
                err.message);
        } else {
            console.log('reverse for ' + a + ': ' +
                JSON.stringify(domains));
        }
    });
    });
});
```

dns.lookup(domain, family=null, callback)

Resolves a domain (e.g. `'google.com'`) into the first found A (IPv4) or AAAA (IPv6) record.

The callback has arguments `(err, address, family)`. The `address` argument is a string representation of a IP v4 or v6 address. The `family` argument is either the integer 4 or 6 and denotes the family of `address` (not neccessarily the value initially passed to `lookup`).

dns.resolve(domain, rrtype='A', callback)

Resolves a domain (e.g. `'google.com'`) into an array of the record types specified by rrtype. Valid rrtypes are `A` (IPV4 addresses), `AAAA` (IPV6 addresses), `MX` (mail exchange records), `TXT` (text records), `SRV` (SRV records), and `PTR` (used for reverse IP lookups).

The callback has arguments `(err, addresses)`. The type of each item in `addresses` is determined by the record type, and described in the documentation for the corresponding lookup methods below.

On error, `err` would be an instanceof `Error` object, where `err.errno` is one of the error codes listed below and `err.message` is a string describing the error in English.

dns.resolve4(domain, callback)

The same as `dns.resolve()`, but only for IPv4 queries (`A` records). `addresses` is an array of IPv4 addresses (e.g.

`['74.125.79.104', '74.125.79.105', '74.125.79.106']`).

dns.resolve6(domain, callback)

The same as `dns.resolve4()` except for IPv6 queries (an `AAAA` query).

dns.resolveMx(domain, callback)

The same as `dns.resolve()`, but only for mail exchange queries (`MX` records).

`addresses` is an array of MX records, each with a priority and an exchange attribute (e.g.

`[{'priority': 10, 'exchange': 'mx.example.com'},...]`).

dns.resolveTxt(domain, callback)

The same as `dns.resolve()`, but only for text queries (`TXT` records). `addresses` is an array of the text records available for `domain` (e.g., `['v=spf1 ip4:0.0.0.0 ~all']`).

dns.resolveSrv(domain, callback)

The same as `dns.resolve()`, but only for service records (SRV records). `addresses` is an array of the SRV records available for `domain`. Properties of SRV records are priority, weight, port, and name (e.g., `[{'priority': 10, {'weight': 5, 'port': 21223, 'name': 'service.example.com'}, ...]`).

dns.reverse(ip, callback)

Reverse resolves an ip address to an array of domain names.

The callback has arguments `(err, domains)`.

If there an an error, `err` will be non-null and an instanceof the Error object.

Each DNS query can return an error code.

- `dns.TEMPFAIL`: timeout, SERVFAIL or similar.
- `dns.PROTOCOL`: got garbled reply.
- `dns.NXDOMAIN`: domain does not exists.
- `dns.NODATA`: domain exists but no data of reqd type.
- `dns.NOMEM`: out of memory while processing.
- `dns.BADQUERY`: the query is malformed.

dgram

Datagram sockets are available through `require('dgram')`. Datagrams are most commonly handled as IP/UDP messages, but they can also be used over Unix domain sockets.

Event: 'message'
`function (msg, rinfo) { }`

Emitted when a new datagram is available on a socket. `msg` is a `Buffer` and `rinfo` is an object with the sender's address information and the number of bytes in the datagram.

Event: 'listening'
`function () { }`

Emitted when a socket starts listening for datagrams. This happens as soon as UDP sockets are created. Unix domain sockets do not start listening until calling `bind()` on them.

Event: 'close'
`function () { }`

Emitted when a socket is closed with `close()`. No new `message` events will be emitted on this socket.

dgram.createSocket(type, [callback])

Creates a datagram socket of the specified types. Valid types are: `udp4`, `udp6`, and `unix_dgram`.

Takes an optional callback which is added as a listener for `message` events.

dgram.send(buf, offset, length, path, [callback])

For Unix domain datagram sockets, the destination address is a pathname in the filesystem. An optional callback may be supplied that is invoked after the `sendto` call is completed by the OS. It is not safe to re-use `buf` until the callback is invoked. Note that unless the socket is bound to a pathname with `bind()` there is no way to receive messages on this socket.

Example of sending a message to syslogd on OSX via Unix domain socket `/var/run/syslog`:

```
var dgram = require('dgram');
var message = new Buffer("A message to log.");
var client = dgram.createSocket("unix_dgram");
client.send(message, 0, message.length, "/var/run/syslog",
  function (err, bytes) {
    if (err) {
      throw err;
    }
    console.log("Wrote " + bytes + " bytes to socket.");
});
```

dgram.send(buf, offset, length, port, address, [callback])

For UDP sockets, the destination port and IP address must be specified. A string may be supplied for the `address` parameter, and it will be resolved with DNS. An optional callback may be specified to detect any DNS errors and when `buf` may be re-used. Note that DNS lookups will delay the time that a send takes place, at least until the next tick. The only way to know for sure that a send has taken place is to use the callback.

Example of sending a UDP packet to a random port on `localhost`;

```
var dgram = require('dgram');
var message = new Buffer("Some bytes");
var client = dgram.createSocket("udp4");
client.send(message, 0, message.length, 41234, "localhost");
client.close();
```

dgram.bind(path)

For Unix domain datagram sockets, start listening for incoming datagrams on a socket specified by `path`. Note that clients may `send()` without `bind()`, but no datagrams will be received without a `bind()`.

Example of a Unix domain datagram server that echoes back all messages it receives:

```
var dgram = require("dgram");
var serverPath = "/tmp/dgram_server_sock";
var server = dgram.createSocket("unix_dgram");

server.on("message", function (msg, rinfo) {
  console.log("got: " + msg + " from " + rinfo.address);
  server.send(msg, 0, msg.length, rinfo.address);
});

server.on("listening", function () {
  console.log("server listening " + server.address().address);
})

server.bind(serverPath);
```

Example of a Unix domain datagram client that talks to this server:

```
var dgram = require("dgram");
var serverPath = "/tmp/dgram_server_sock";
var clientPath = "/tmp/dgram_client_sock";

var message = new Buffer("A message at " + (new Date()));

var client = dgram.createSocket("unix_dgram");

client.on("message", function (msg, rinfo) {
  console.log("got: " + msg + " from " + rinfo.address);
});

client.on("listening", function () {
  console.log("client listening " + client.address().address);
  client.send(message, 0, message.length, serverPath);
});

client.bind(clientPath);
```

dgram.bind(port, [address])
For UDP sockets, listen for datagrams on a named `port` and optional `address`. If `address` is not specified, the OS will try to listen on all addresses.

Example of a UDP server listening on port 41234:

```
var dgram = require("dgram");

var server = dgram.createSocket("udp4");
var messageToSend = new Buffer("A message to send");
```

```
server.on("message", function (msg, rinfo) {
  console.log("server got: " + msg + " from " +
    rinfo.address + ":" + rinfo.port);
});

server.on("listening", function () {
  var address = server.address();
  console.log("server listening " +
      address.address + ":" + address.port);
});

server.bind(41234);
// server listening 0.0.0.0:41234
```

dgram.close()

Close the underlying socket and stop listening for data on it. UDP sockets automatically listen for messages, even if they did not call `bind()`.

dgram.address()

Returns an object containing the address information for a socket. For UDP sockets, this object will contain `address` and `port`. For Unix domain sockets, it will contain only `address`.

dgram.setBroadcast(flag)

Sets or clears the `SO_BROADCAST` socket option. When this option is set, UDP packets may be sent to a local interface's broadcast address.

dgram.setTTL(ttl)

Sets the `IP_TTL` socket option. TTL stands for "Time to Live," but in this context it specifies the number of IP hops that a packet is allowed to go through. Each router or gateway that forwards a packet decrements the TTL. If the TTL is decremented to 0 by a router, it will not be forwarded. Changing TTL values is typically done for network probes or when multicasting.

The argument to `setTTL()` is a number of hops between 1 and 255. The default on most systems is 64.

Assert

This module is used for writing unit tests for your applications, you can access it with `require('assert')`.

assert.fail(actual, expected, message, operator)

Tests if `actual` is equal to `expected` using the operator provided.

assert.ok(value, [message])

Tests if value is a `true` value, it is equivalent to `assert.equal(true, value, message);`

assert.equal(actual, expected, [message])

Tests shallow, coercive equality with the equal comparison operator (`==`).

assert.notEqual(actual, expected, [message])

Tests shallow, coercive non-equality with the not equal comparison operator (`!=`).

assert.deepEqual(actual, expected, [message])

Tests for deep equality.

assert.notDeepEqual(actual, expected, [message])

Tests for any deep inequality.

assert.strictEqual(actual, expected, [message])

Tests strict equality, as determined by the strict equality operator (`===`)

assert.notStrictEqual(actual, expected, [message])

Tests strict non-equality, as determined by the strict not equal operator (`!==`)

assert.throws(block, [error], [message])

Expects `block` to throw an error.

assert.doesNotThrow(block, [error], [message])

Expects `block` not to throw an error.

assert.ifError(value)

Tests if value is not a false value, throws if it is a true value. Useful when testing the first argument, `error` in callbacks.

Path

This module contains utilities for dealing with file paths. Use `require('path')` to use it. It provides the following methods:

path.join([path1], [path2], [...])
Join all arguments together and resolve the resulting path.

Example:

```
node> require('path').join(
...     '/foo', 'bar', 'baz/asdf', 'quux', '..')
'/foo/bar/baz/asdf'
```

path.normalizeArray(arr)
Normalize an array of path parts, taking care of `'..'` and `'.'` parts.

Example:

```
path.normalizeArray(['',
   'foo', 'bar', 'baz', 'asdf', 'quux', '..'])
// returns
[ '', 'foo', 'bar', 'baz', 'asdf' ]
```

path.normalize(p)

Normalize a string path, taking care of `'..'` and `'.'` parts.

Example:

```
path.normalize('/foo/bar/baz/asdf/quux/..')
// returns
'/foo/bar/baz/asdf'
```

path.dirname(p)

Return the directory name of a path. Similar to the Unix `dirname` command.

Example:

```
path.dirname('/foo/bar/baz/asdf/quux')
// returns
'/foo/bar/baz/asdf'
```

path.basename(p, [ext])

Return the last portion of a path. Similar to the Unix `basename` command.

Example:

```
path.basename('/foo/bar/baz/asdf/quux.html')
// returns
'quux.html'

path.basename('/foo/bar/baz/asdf/quux.html', '.html')
// returns
'quux'
```

path.extname(p)

Return the extension of the path. Everything after the last '.' in the last portion of the path.
If there is no '.' in the last portion of the path or the only '.' is the first character, then it
returns an empty string. Examples:

```
path.extname('index.html')
// returns
'.html'

path.extname('index')
// returns
''
```

path.exists(p, [callback])

Test whether or not the given path exists. Then, call the `callback` argument with either true or false. Example:

```
path.exists('/etc/passwd', function (exists) {
  sys.debug(exists ? "it's there" : "no passwd!");
});
```

URL

This module has utilities for URL resolution and parsing. Call `require('url')` to use it.

Parsed URL objects have some or all of the following fields, depending on whether or not they exist in the URL string. Any parts that are not in the URL string will not be in the parsed object. Examples are shown for the URL

`'http://user:pass@host.com:8080/p/a/t/h?query=string#hash'`

- `href`

The full URL that was originally parsed. Example:
`'http://user:pass@host.com:8080/p/a/t/h?query=string#hash'`

- `protocol`

The request protocol. Example: `'http:'`

- `host`

The full host portion of the URL, including port and authentication information. Example:
`'user:pass@host.com:8080'`

- `auth`

The authentication information portion of a URL. Example: `'user:pass'`

- `hostname`

Just the hostname portion of the host. Example: `'host.com'`

- `port`

The port number portion of the host. Example: `'8080'`

- `pathname`

The path section of the URL, that comes after the host and before the query, including the

initial slash if present. Example: `'/p/a/t/h'`

- **search**

The 'query string' portion of the URL, including the leading question mark. Example: `'?query=string'`

- **query**

Either the 'params' portion of the query string, or a querystring-parsed object. Example: `'query=string'` or `{'query':'string'}`

- **hash**

The 'fragment' portion of the URL including the pound-sign. Example: `'#hash'`

The following methods are provided by the URL module:

url.parse(urlStr, parseQueryString=false)
Take a URL string, and return an object. Pass `true` as the second argument to also parse the query string using the `querystring` module.

url.format(urlObj)
Take a parsed URL object, and return a formatted URL string.

url.resolve(from, to)
Take a base URL, and a href URL, and resolve them as a browser would for an anchor tag.

Query String

This module provides utilities for dealing with query strings. It provides the following methods:

querystring.stringify(obj, sep='&', eq='=')
Serialize an object to a query string. Optionally override the default separator and assignment characters.

Example:

```
querystring.stringify({foo: 'bar'})
// returns
'foo=bar'

querystring.stringify({foo: 'bar', baz: 'bob'}, ';', ':')
// returns
'foo:bar;baz:bob'
```

querystring.parse(str, sep='&', eq='=')

Deserialize a query string to an object. Optionally override the default separator and assignment characters.

Example:

```
querystring.parse('a=b&b=c')
// returns
{ 'a': 'b'
, 'b': 'c'
}
```

querystring.escape

The escape function used by `querystring.stringify`, provided so that it could be overridden if necessary.

querystring.unescape

The unescape function used by `querystring.parse`, provided so that it could be overridden if necessary.

REPL

A Read-Eval-Print-Loop (REPL) is available both as a standalone program and easily includable in other programs. REPL provides a way to interactively run JavaScript and see the results. It can be used for debugging, testing, or just trying things out.

By executing `node` without any arguments from the command-line you will be dropped into the REPL. It has simplistic emacs line-editing.

```
mjr:~$ node
Type '.help' for options.
node> a = [ 1, 2, 3];
[ 1, 2, 3 ]
node> a.forEach(function (v) {
...     console.log(v);
...     });
1
2
3
```

For advanced line-editors, start node with the environmental variable `NODE_NO_READLINE=1`. This will start the REPL in canonical terminal settings which will allow you to use with `rlwrap`.

For example, you could add this to your bashrc file:

```
alias node="env NODE_NO_READLINE=1 rlwrap node"
```

repl.start(prompt='node> ', stream=process.openStdin())

Starts a REPL with `prompt` as the prompt and `stream` for all I/O. `prompt` is optional and defaults to `node>` . `stream` is optional and defaults to `process.openStdin()`.

Multiple REPLs may be started against the same running instance of node. Each will share the same global object but will have unique I/O.

Here is an example that starts a REPL on stdin, a Unix socket, and a TCP socket:

```
var net = require("net"),
    repl = require("repl");

connections = 0;

repl.start("node via stdin> ");

net.createServer(function (socket) {
  connections += 1;
  repl.start("node via Unix socket> ", socket);
}).listen("/tmp/node-repl-sock");

net.createServer(function (socket) {
  connections += 1;
  repl.start("node via TCP socket> ", socket);
}).listen(5001);
```

Running this program from the command line will start a REPL on stdin. Other REPL clients may connect through the Unix socket or TCP socket. `telnet` is useful for connecting to TCP sockets, and `socat` can be used to connect to both Unix and TCP sockets.

By starting a REPL from a Unix socket-based server instead of stdin, you can connect to a long-running node process without restarting it.

REPL Features

Inside the REPL, Control+D will exit. Multi-line expressions can be input.

The special variable _ (underscore) contains the result of the last expression.

```
node> [ "a", "b", "c" ]
[ 'a', 'b', 'c' ]
node> _.length
3
node> _ += 1
4
```

The REPL provides access to any variables in the global scope. You can expose a variable to the REPL explicitly by assigning it to the `context` object associated with each `REPLServer`. For example:

```
// repl_test.js
var repl = require("repl"),
    msg = "message";

repl.start().context.m = msg;
```

Things in the `context` object appear as local within the REPL:

```
mjr:~$ node repl_test.js
node> m
'message'
```

There are a few special REPL commands:

- `.break` - While inputting a multi-line expression, sometimes you get lost or just don't care about completing it. `.break` will start over.

- `.clear` - Resets the `context` object to an empty object and clears any multi-line expression.

- `.exit` - Close the I/O stream, which will cause the REPL to exit.

- `.help` - Show this list of special commands.

Modules

Node uses the CommonJS module system.

Node has a simple module loading system. In Node, files and modules are in one-to-one correspondence. As an example, `foo.js` loads the module `circle.js` in the same directory.

The contents of `foo.js`:

```
var circle = require('./circle');
console.log( 'The area of a circle of radius 4 is '
            + circle.area(4));
```

The contents of `circle.js`:

```
var PI = 3.14;

exports.area = function (r) {
  return PI * r * r;
};

exports.circumference = function (r) {
  return 2 * PI * r;
```

```
    };
```

The module `circle.js` has exported the functions `area()` and `circumference()`. To export an object, add to the special `exports` object. (Alternatively, one can use `this` instead of `exports`.) Variables local to the module will be private. In this example the variable `PI` is private to `circle.js`. The function `puts()` comes from the module `'sys'`, which is a built-in module. Modules which are not prefixed by `'./'` are built-in module--more about this later.

A module prefixed with `'./'` is relative to the file calling `require()`. That is, `circle.js` must be in the same directory as `foo.js` for `require('./circle')` to find it.

Without the leading `'./'`, like `require('assert')` the module is searched for in the `require.paths` array. `require.paths` on my system looks like this:

```
[ '/home/ryan/.node_libraries' ]
```

That is, when `require('assert')` is called Node looks for:

- 1: `/home/ryan/.node_libraries/assert.js`
- 2: `/home/ryan/.node_libraries/assert.node`
- 3: `/home/ryan/.node_libraries/assert/index.js`
- 4: `/home/ryan/.node_libraries/assert/index.node`

interrupting once a file is found. Files ending in `'.node'` are binary Addon Modules; see 'Addons' below. `'index.js'` allows one to package a module as a directory.

`require.paths` can be modified at runtime by simply unshifting new paths onto it, or at startup with the `NODE_PATH` environmental variable (which should be a list of paths, colon separated).

Addons

Addons are dynamically linked shared objects. They can provide glue to C and C++ libraries. The API (at the moment) is rather complex, involving knowledge of several libraries:

- V8 JavaScript, a C++ library. Used for interfacing with JavaScript: creating objects, calling functions, etc. Documented mostly in the `v8.h` header file (`deps/v8/include/v8.h` in the Node source tree).

- libev, C event loop library. Anytime one needs to wait for a file descriptor to become readable, wait for a timer, or wait for a signal to received one will need to interface with libev. That is, if you perform any I/O, libev will need to be used. Node uses the `EV_DEFAULT` event loop. Documentation can be found http:/cvs.schmorp.de/libev/ev.html[here].

- libeio, C thread pool library. Used to execute blocking POSIX system calls

asynchronously. Mostly wrappers already exist for such calls, in `src/file.cc` so you will probably not need to use it. If you do need it, look at the header file `deps/libeio/eio.h`.

- Internal Node libraries. Most importantly is the `node::ObjectWrap` class which you will likely want to derive from.

- Others. Look in `deps/` for what else is available.

Node statically compiles all its dependencies into the executable. When compiling your module, you don't need to worry about linking to any of these libraries.

To get started let's make a small Addon which does the following except in C++:

```
exports.hello = 'world';
```

To get started we create a file `hello.cc`:

```
#include <v8.h>

using namespace v8;

extern "C" void
init (Handle<Object> target)
{
  HandleScope scope;
  target->Set(String::New("hello"), String::New("World"));
}
```

This source code needs to be built into `hello.node`, the binary Addon. To do this we create a file called `wscript` which is python code and looks like this:

```
srcdir = '.'
blddir = 'build'
VERSION = '0.0.1'

def set_options(opt):
  opt.tool_options('compiler_cxx')

def configure(conf):
  conf.check_tool('compiler_cxx')
  conf.check_tool('node_addon')

def build(bld):
  obj = bld.new_task_gen('cxx', 'shlib', 'node_addon')
  obj.target = 'hello'
  obj.source = 'hello.cc'
```

Running `node-waf configure build` will create a file `build/default/hello.node` which is our Addon.

`node-waf` is just http://code.google.com/p/waf/[WAF], the python-based build system. `node-waf` is provided for the ease of users.

All Node addons must export a function called `init` with this signature:

```
extern 'C' void init (Handle<Object> target)
```

For the moment, that is all the documentation on addons. Please see http://github.com/ry/node_postgres for a real example.

Appendix - Third Party Modules

There are many third party modules for Node. At the time of writing, August 2010, the master repository of modules is http://github.com/ry/node/wiki/modules[the wiki page].

This appendix is intended as a SMALL guide to new-comers to help them quickly find what are considered to be quality modules. It is not intended to be a complete list. There may be better more complete modules found elsewhere.

- Module Installer: npm

- HTTP Middleware: Connect

- Web Framework: Express

- Web Sockets: Socket.IO

- HTML Parsing: HTML5

- mDNS/Zeroconf/Bonjour

- RabbitMQ, AMQP

- mysql

- Serialization: msgpack

- Scraping: Apricot

- Debugger: ndb is a CLI debugger inspector is a web based tool.

- pcap binding

- ncurses

- Testing/TDD/BDD: vows, expresso, mjsunit.runner

Patches to this list are welcome.